

11. Evaluating Binary Classification

Bruce E. Shapiro

Getting Started in Machine Learning

Copyright (c) 2019. May not be distributed in any form without written permission.

Last revised: March 27, 2019

Image you are a Radar Operator During WW II



Image from Yank Magazine, Oct. 5, 1945.

The “blips” you don’t recognize (and classify as enemy aircraft) may be

- Enemy aircraft
 - ▶ We call these TRUE POSITIVES
- False alarms (birds, weather, noise, etc.)
 - ▶ We call these FALSE POSITIVES

Image you are a Radar Operator During WW II



Image from Yank Magazine, Oct. 5, 1945.

The “blips” you don’t recognize (and classify as enemy aircraft) may be

- Enemy aircraft
 - ▶ We call these TRUE POSITIVES
- False alarms (birds, weather, noise, etc.)
 - ▶ We call these FALSE POSITIVES

Other blips whose radar signatures you think you recognize could be

- Friendly aircraft, birds, weather, etc., correctly identified.
 - ▶ We call these TRUE NEGATIVES
- Enemy aircraft that not identified as threats
 - ▶ We call these FALSE NEGATIVES

Confusion Matrix

- count **true positive** (TP) if $y = 1$ and $y_p = 1$
- count **false positive** (FP) if $y = 0$ and $y_p = 1$
- count **true negative** (TN) if $y = 0$ and $y_p = 0$
- count **false negative** (FN) if $y = 1$ and $y_p = 0$

		Predicted Category y_p	
		0	1
Actual Category y	0	True Negative TN	False Positive FP
	1	False Negative FN	True Positive TP

Measures of Success

- Let $N = TN + FP$ = total number of $y = 0$'s (negative examples)
- Let $P = TP + FN$ = total number of $y = 1$'s (positive examples)

- **True Positive Rate, Recall, Sensitivity**

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN}$$

- **specificity** = $\frac{TN}{N} = \frac{TN}{FP + TN}$

- **precision** = $\frac{TP}{TP + FP}$

- **accuracy** = $\frac{TP + TN}{P + N}$

- **False Positive Rate, Fallout**

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN} \\ = 1 - \text{specificity}$$

- **error rate** = $1 - \text{accuracy}$

Pythonic Calculation of TP, TN, FP, FN

- Python stores **booleans** as **1** for **True** and **0** for **False**
- The expression **True and False** evaluates to **False** and is stored as a **0**
- The expression **False and False** evaluates to **False** and is stored as a **0**
- The expression **True and True** evaluates to **True** and is stored as a **1**
- Similarly, the expression **X==1 and Y==0** will evaluate to **True** (i.e., **1**) only when **X** is **1** and **Y** is **0**, and will evaluate to **False** (i.e., **0**) otherwise

To count the TP, TN, FP, FN manually:

```
OPS=list(zip(Y,YP))
TP = sum([(OB==1) and (PR==1) for OB,PR in OPS])
TN = sum([(OB==0) and (PR==0) for OB,PR in OPS])
FP = sum([(PR==1) and (OB==0) for OB,PR in OPS])
FN = sum([(PR==0) and (OB==1) for OB,PR in OPS])
```

here **YP** is the output of a **predict** method and **Y** is the array of exemplars (such as **YTEST**)

To count the TP, TN, FP, FN manually:

```
OPS=list(zip(Y,YP))
TP = sum([(OB==1) and (PR==1) for OB,PR in OPS])
TN = sum([(OB==0) and (PR==0) for OB,PR in OPS])
FP = sum([(PR==1) and (OB==0) for OB,PR in OPS])
FN = sum([(PR==0) and (OB==1) for OB,PR in OPS])
```

here **YP** is the output of a **predict** method and **Y** is the array of exemplars (such as **YTEST**)

```
confusionMatrix=np.array([[TN,FP],[FN,TP]])
```


To count the TP, TN, FP, FN manually:

```
OPS=list(zip(Y,YP))
TP = sum([(OB==1) and (PR==1) for OB,PR in OPS])
TN = sum([(OB==0) and (PR==0) for OB,PR in OPS])
FP = sum([(PR==1) and (OB==0) for OB,PR in OPS])
FN = sum([(PR==0) and (OB==1) for OB,PR in OPS])
```

here **YP** is the output of a **predict** method and **Y** is the array of exemplars (such as **YTEST**)

```
confusionMatrix=np.array([[TN,FP],[FN,TP]])
```

```
P = sum(OBS)
N = len(OBS)-P
TPR = TP/P # True Positive Rate or Recall
TNR = TN/N
Specificity = TN/(FP+TN)
Accuracy = (TP+TN)/(P+N)
Precision = TP/(TP+FP)
```

Equivalent Metrics in Sklearn

```
sklearn.metrics.confusion_matrix(YTEST, YP)  
sklearn.metrics.recall_score(YTEST, YP)  
sklearn.metrics.accuracy_score(YTEST, YP)  
sklearn.metrics.precision_score(YTEST, YP)
```

ROC Curve

- Plots **true positive rate** as a function of **false positive rate** at different thresholds
- The threshold is varied by sorting the data in decreasing order by probability

ROC Curve

- Plots **true positive rate** as a function of **false positive rate** at different thresholds
- The threshold is varied by sorting the data in decreasing order by probability

```
def ROC(Y, Prob) :  
    if (len(Y) != len(Prob)) :  
        print("Length mismatch")  
        return([])  
  
    P=sum(Y) ; N=len(Y)-P  
  
    if (P<1) or (N<1):  
        print("There must be both postive and negative example")  
        return([])  
  
    # code continued on next page ...
```


ROC (continued)

```
py_pairs = sorted(zip(Prob, Y), reverse=True)

FP=0; TP=0
ROC_CURVE=[]
pprev=float("-inf")

for p,y in py_pairs:
    if p != pprev:
        ROC_CURVE.append([FP/N, TP/P])
        pprev=p
    if y>0:
        TP+=1
    else:
        FP+=1

ROC_CURVE.append([FP/N, TP/P])
return(ROC_CURVE)
```



Code for Plotting ROC Curve With Sklearn

After reading data, and creating test and training set: (this example is for Logistic regression but it works for any method with two classes.)

```
model=LR().fit(XTRAIN,YTRAIN)
probs=model.predict_proba(XTEST)[: ,1]
fpr, tpr, threshold = roc_curve(YTEST,probs)
#
# <-- Insert code for threshold annotations here
#      (see next slide)
#
plt.plot(fpr,tpr, c="blue")
plt.plot([0,1],[0,1],c="green",ls="--")
plt.xlabel("False Positive Rate", fontsize=14)
plt.ylabel("True Positive Rate", fontsize=14)
plt.grid()
```

Code for Annotating ROC Curve with Thresholds

```
from scipy import interpolate
# ----- Create an interpolating functions
x=threshold          # domain of f
y=np.vstack([fpr, tpr]) # range of f
f=interpolate.interpld(x,y)
# ----- Interpolate at interesting threshold points
tvals=np.arange(.2, .9, .1)
xyvals=f(tvals)
xvals,yvals=xyvals
# ----- Add scatter plot to the figure
for t,x,y in zip(tvals,xvals,yvals):
    plt.annotate(str(round(t,1)), (x,y),
        xytext=(x+.1,y-.1), fontsize=14,
        arrowprops={"width":1.0, "facecolor":"black",
            "headwidth":6})
```


Citations

- Radar image from Yank Magazine, Oct 1945, Copyright unclear, posted at http://www.oldmagazinearticles.com/WW2_radar_history_article_development_of_radar_in_world_war_two#.XHPr-HWQE5k
- UCI MPG data set
<https://archive.ics.uci.edu/ml/datasets/auto+mpg>;
relevant paper: Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.
- Dua, D. and Karra Taniskidou, E. (2017). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science.
- ROC Curve Algorithm: Tom Fawcett (2006) An Introduction to ROC analysis, Pattern Recognition Letters, 27:861-874.