

# 19. Tree Ensembles

Bruce E. Shapiro

## Getting Started in Machine Learning

Copyright (c) 2019. May not be distributed in any form without  
written permission.

Last revised: March 31, 2019

# Ensembles of Trees

- Compare the results of many decision trees.
- `sklearn.ensemble` library
  - ▶ **Boosting**. Computes error; iterates and finds weighted average.
    - **Adaptive Boosting** (Adaboost) - modifies weights of sample
    - **Gradient Boosting** - minimizes errors by gradient descent
  - ▶ **Bagging** or **bootstrap aggregation**. Generates multiple training sets and forms weighted average.
  - ▶ **Random Forests**. Uses a subset of features in each subset.
  - ▶ **Isolation Forests**. Variant of random forest that computes path length, thereby identifying outliers (shorter paths).
  - ▶ ...



# Wine data file

fixed acid- ity	volatile acidity	citric acid	residual sugar	chlor- ides	free sulfur dioxide	total sulfur dioxide	density	pH	sufph-	alcohol	quality
7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	6
6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	6
8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	6
		.									
		.									
		.									
		(etc)									

## Read File, Create x, y Arrays

```
import pandas as pd
import numpy as np
white=pd.read_csv("https://archive.ics.uci.edu/ml/
    machine-learning-databases/wine-quality/
    winequality-white.csv", sep=";")
Y=np.array(white["quality"])
X=np.array(white)[:,:-1]
```

# Forests and Trees: Approach

- for each forest size  $n = 2, \dots$ :
  - ▶ repeat  $N$  times:
    - perform a train/test split
    - calculate the classification error
  - ▶ calculate:
    - $\mu$  = mean of  $N$  classification errors
    - $\sigma$  = standard deviation of  $N$  classification errors
- plot  $\mu \pm \sigma$  vs forest size  $n$  (number of trees)

## Code to Do `nruns` Train/Test splits

- let `clf` is a classifier from `sklearn`:
- Run `nruns` train/test split

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
def evaluate_classifier(clf, X, Y, nruns):
    errs=[]
    for j in range(nruns):
        XTRAIN, XTEST, YTRAIN, YTEST=\
            train_test_split(X,Y)
        clf.fit(XTRAIN, YTRAIN)
        YP=clf.predict(XTEST)
        errs.append(1-accuracy_score(YTEST, YP))
    return(np.mean(errs), np.std(errs))
```

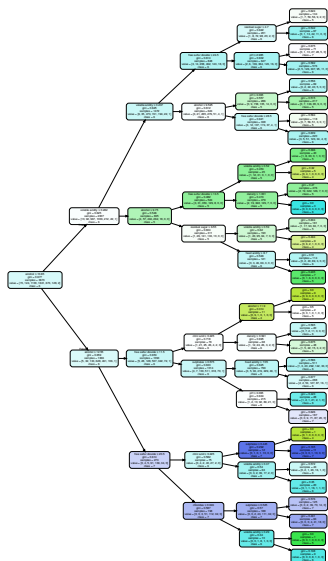
- Return value is  $(\mu, \sigma)$

## Baselines: Single Tree (no ensemble)

```
DT5=DecisionTreeClassifier(max_depth=5)
DTT5mu,DTT5sig=evaluate_classifier(DT5,X,Y,25)
DT=DecisionTreeClassifier(max_depth=None)
DTFullmu, DTFullsig=evaluate_classifier(DT,X,Y,25)
print("Max Depth = 5: ", round(DTT5mu,3),
      round(DTT5sig,3))
print("Max Depth = No Limit: ", round(DTFullmu,3),
      round(DTFullsig,3))
```

```
Max Depth = 5:  0.47 0.015
Max Depth = No Limit:  0.407 0.015
```

# Depth-5 Decision Tree





## Decision Tree - No Depth Limit

Full tree at <https://github.com/biomathman/machine-learning-in-python-book/blob/master/extras/white-wine-class.pdf>

# Evaluation of Classifier vs. Forest Size

- Let **clf** be an sklearn ensemble classifier
- Let **forestsizes** be a list of numbers of trees in each forest
- Let **repeats** be the # of train/test repeats for each forest

```
def generate_forests(clf, X, Y, repeats, forestsizes):  
    results=[]  
    for ntrees in forestsizes:  
        classifier=clf(n_estimators=ntrees)  
        mu,sigma=evaluate_classifier(classifier,X,Y,repeats)  
        results.append([mu,sigma])  
        print(ntrees,mu,sigma)  
    mus=np.array(results).T[0]  
    sigs=np.array(results).T[1]  
    return(mus,sigs)
```

- Return value is tuple of lists ( $[\mu_1, \mu_2, \dots], [\sigma_1, \sigma_2, \dots]$ )

# Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier
forest_sizes=[2,3,4,5, 7,10,20,30,40,50,70,100,300,500]
BoostMeans, BoostSigmas=\
    generate_forests(GradientBoostingClassifier,X,Y,100,\
        forest_sizes)
```

```
2 0.47875102040816325 0.013943399597534543
3 0.47075102040816325 0.013394749292965804
4 0.4664326530612245 0.013105933683232481
5 0.46349387755102034 0.01186994917685235
7 0.4559428571428571 0.014189744112709474
10 0.4553714285714287 0.012687014353050224
20 0.44059591836734696 0.011396956271079483
30 0.4323918367346939 0.01581731609807552
40 0.42953469387755105 0.011709309901058097
50 0.42402448979591834 0.014190589416523132
70 0.4157061224489796 0.014030428050738999
100 0.4097795918367347 0.013007020178639313
300 0.3734857142857142 0.012996513164524244
500 0.36040000000000005 0.012070156804161556
```

# Bagging

```
from sklearn.ensemble import BaggingClassifier
BAGmu, BAGsig=\
    generate_forests(BaggingClassifier,X,Y,100,forest_sizes)
```

```
2 0.44858775510204085 0.0144558867383183
3 0.42013877551020407 0.013197278631040045
4 0.3969551020408163 0.014616966515689327
5 0.38026122448979593 0.01435592310209574
7 0.3678612244897959 0.013513424611524709
10 0.3537795918367347 0.012432725420883669
20 0.34132244897959174 0.011983762441976329
30 0.33408979591836735 0.012010613104889854
40 0.33131428571428573 0.011242816948816986
50 0.3295183673469388 0.011165616287904261
70 0.3303591836734694 0.012516016228180335
100 0.32861224489795915 0.011283259872606404
300 0.32581224489795924 0.01089866542082189
500 0.32310204081632654 0.01112842004111149
```

# Random Forests

```
from sklearn.ensemble import RandomForestClassifier
RFMeans, RFSigmas=generate_forests(RandomForestClassifier,
    X,Y,100,forest_sizes)
```

```
2 0.45168979591836733 0.014513119486674303
3 0.4197795918367347 0.015092928304275511
4 0.3940571428571429 0.01423224044397006
5 0.3770040816326531 0.012461247885965058
7 0.3638448979591837 0.01528397635960942
10 0.351526530612245 0.012808150258184626
20 0.3343755102040816 0.013125807168611453
30 0.3298857142857143 0.013076977744829866
40 0.3275591836734694 0.012925117364733224
50 0.3237632653061225 0.013077996884039956
70 0.3228734693877551 0.011163980874726634
100 0.3206448979591836 0.010791759547718657
300 0.31915918367346935 0.012994449204209585
500 0.3174204081632653 0.013213258016927827
```

## Relative Feature Importance (1/2)


```
print(RF.feature_importances_)
```

```
[0.0750037  0.09881211 0.08138214 0.08818542  
 0.08488747 0.09527255 0.09047286 0.10277019  
 0.08668786 0.0787098  0.1178159 ]
```

## Relative Feature Importance (2/2)

```
names=np.array(white.columns)[: -1].tolist()
importances=RF.feature_importances_.tolist()
imp=list(zip(importances,names))
imp.sort(reverse=True)
for impt,name in imp:
    print("%20s %5.3f" %(name,impt/max(importances)))
```

```
          alcohol 1.000
          density 0.872
    volatile acidity 0.839
    free sulfur dioxide 0.809
total sulfur dioxide 0.768
      residual sugar 0.749
              pH 0.736
        chlorides 0.721
        citric acid 0.691
        sulphates 0.668
    fixed acidity 0.637
```





## References

- 1 UCI Machine Learning Repository  
<http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science.
- 2 P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. Modeling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553, 2009.